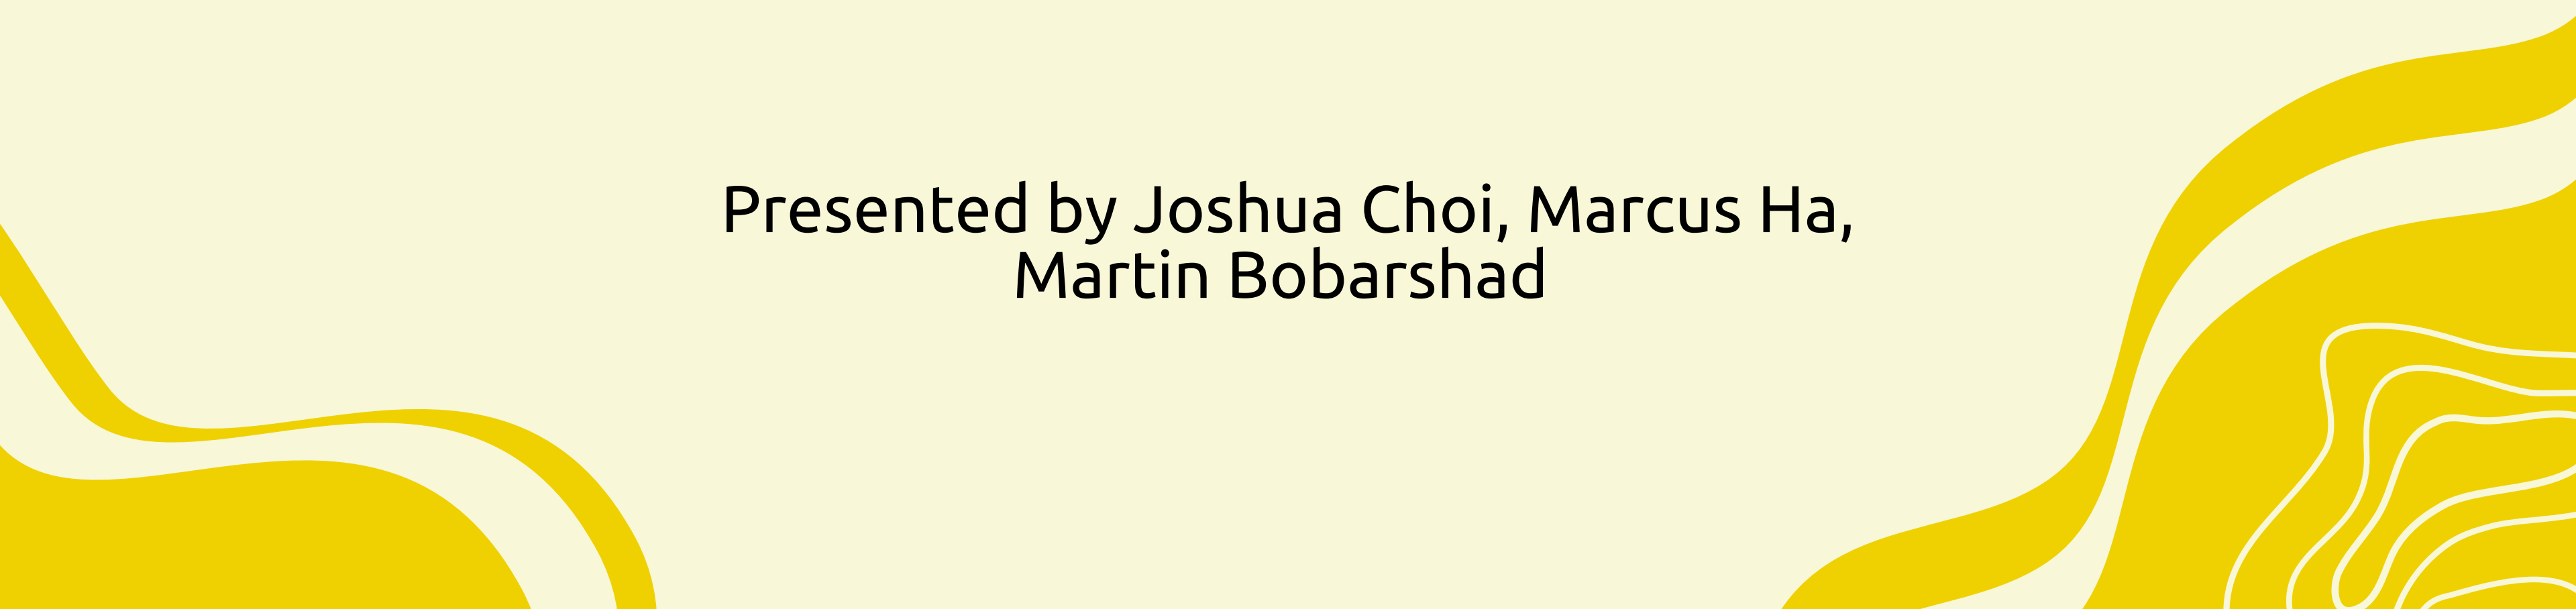



# LANEQUACKER


Presented by Joshua Choi, Marcus Ha,  
Martin Bobarshad



# PROBLEM

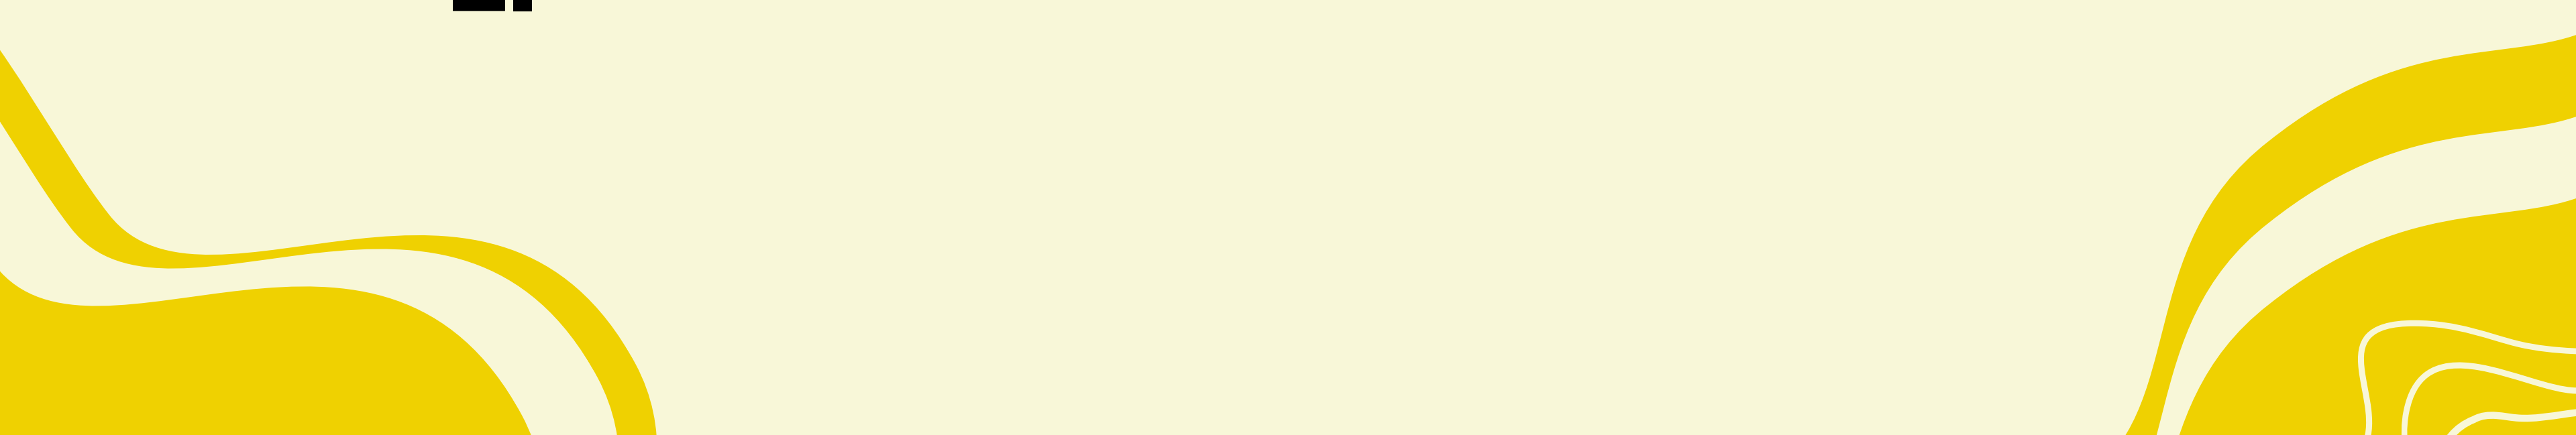


- **It is difficult for ducks to drive!**
  - Continuous and Dynamic Environment
  - Partial Observability (Limited Sensor Input)
  - Complex, Non-discrete Action and Observation Spaces





# PROJECT GOALS

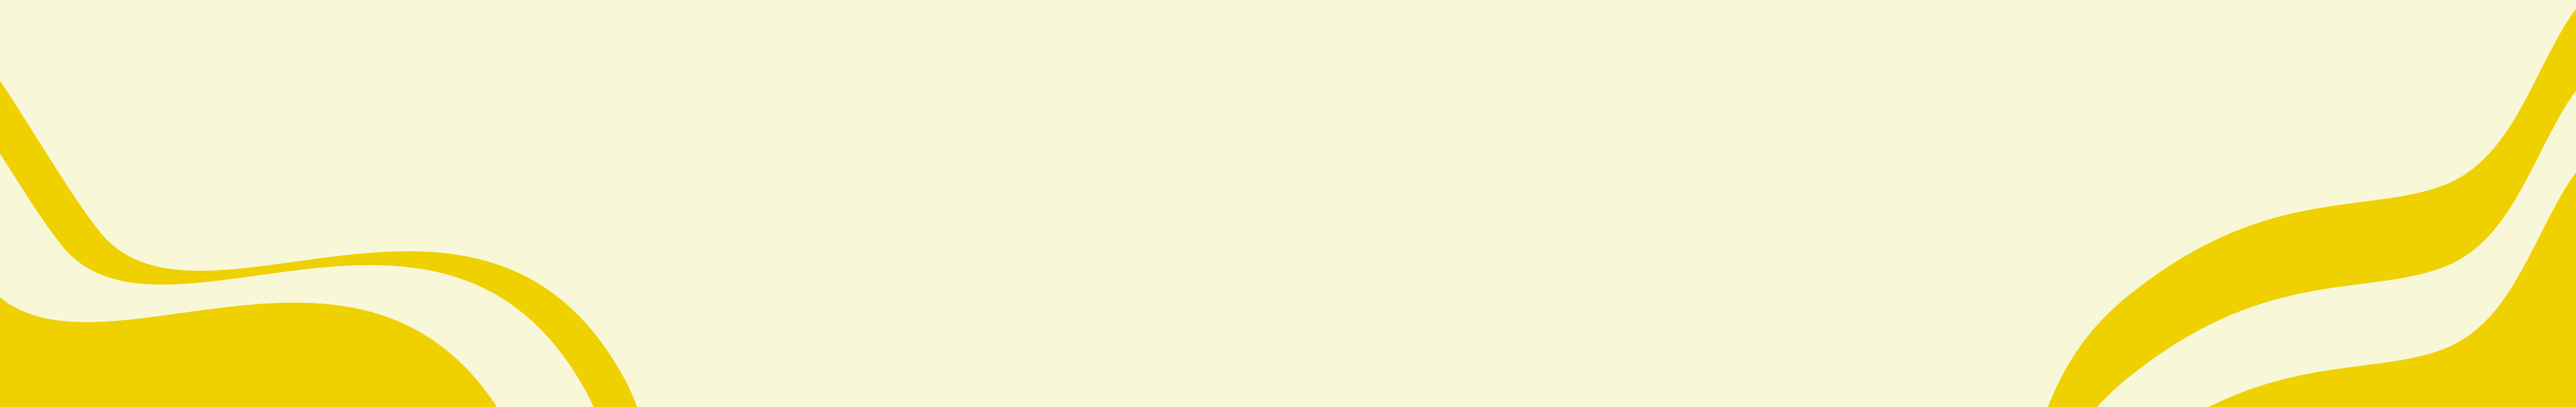
- 1.** Train ducky bot for lane-following, obstacle avoidance, and stopping as needed
  - 2.** Expected outcome: See Duckiebot driving in action!
- 



# APPROACHES

**1. Soft Actor  
Critic (SAC)**


**2. Proximal Policy  
Optimization  
(PPO)**





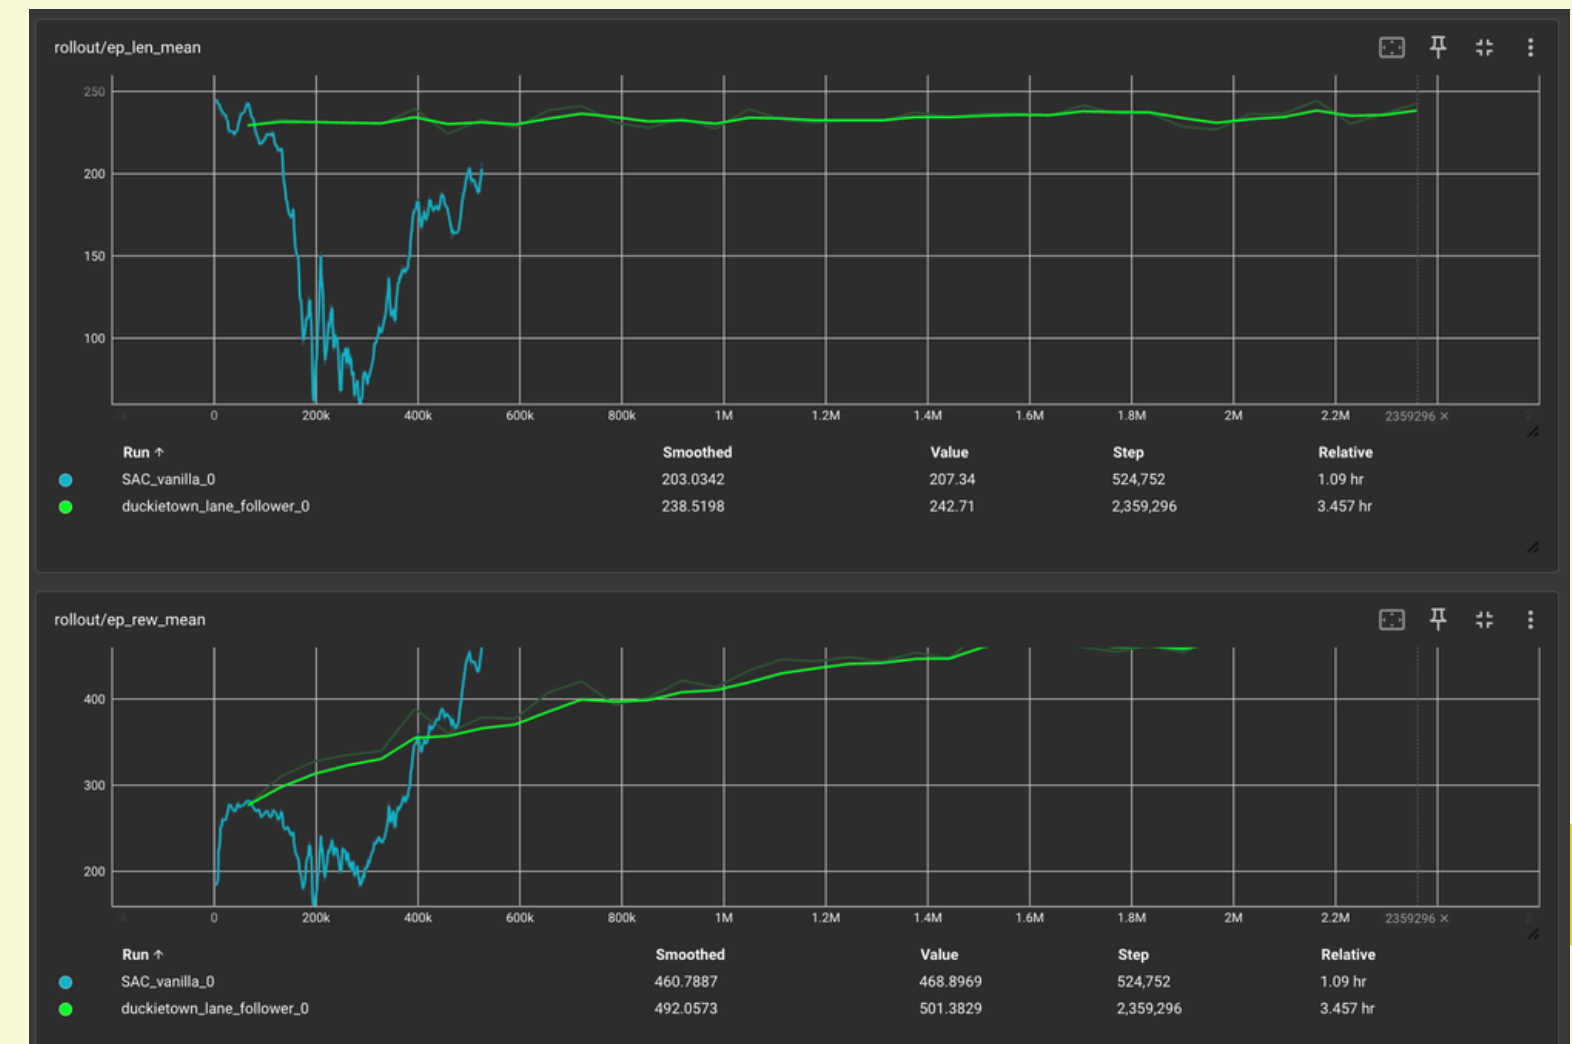
# SOFT ACTOR CRITIC (SAC)

- Off-policy actor-critic RL algorithm
  - Small reward for forward movement
  - Large penalty for collision
  - Small/Large penalty for driving off-lane (depending on distance from center)
- SAC showed less stable training results, making convergence uncertain.
- It required extensive hyperparameter tuning.



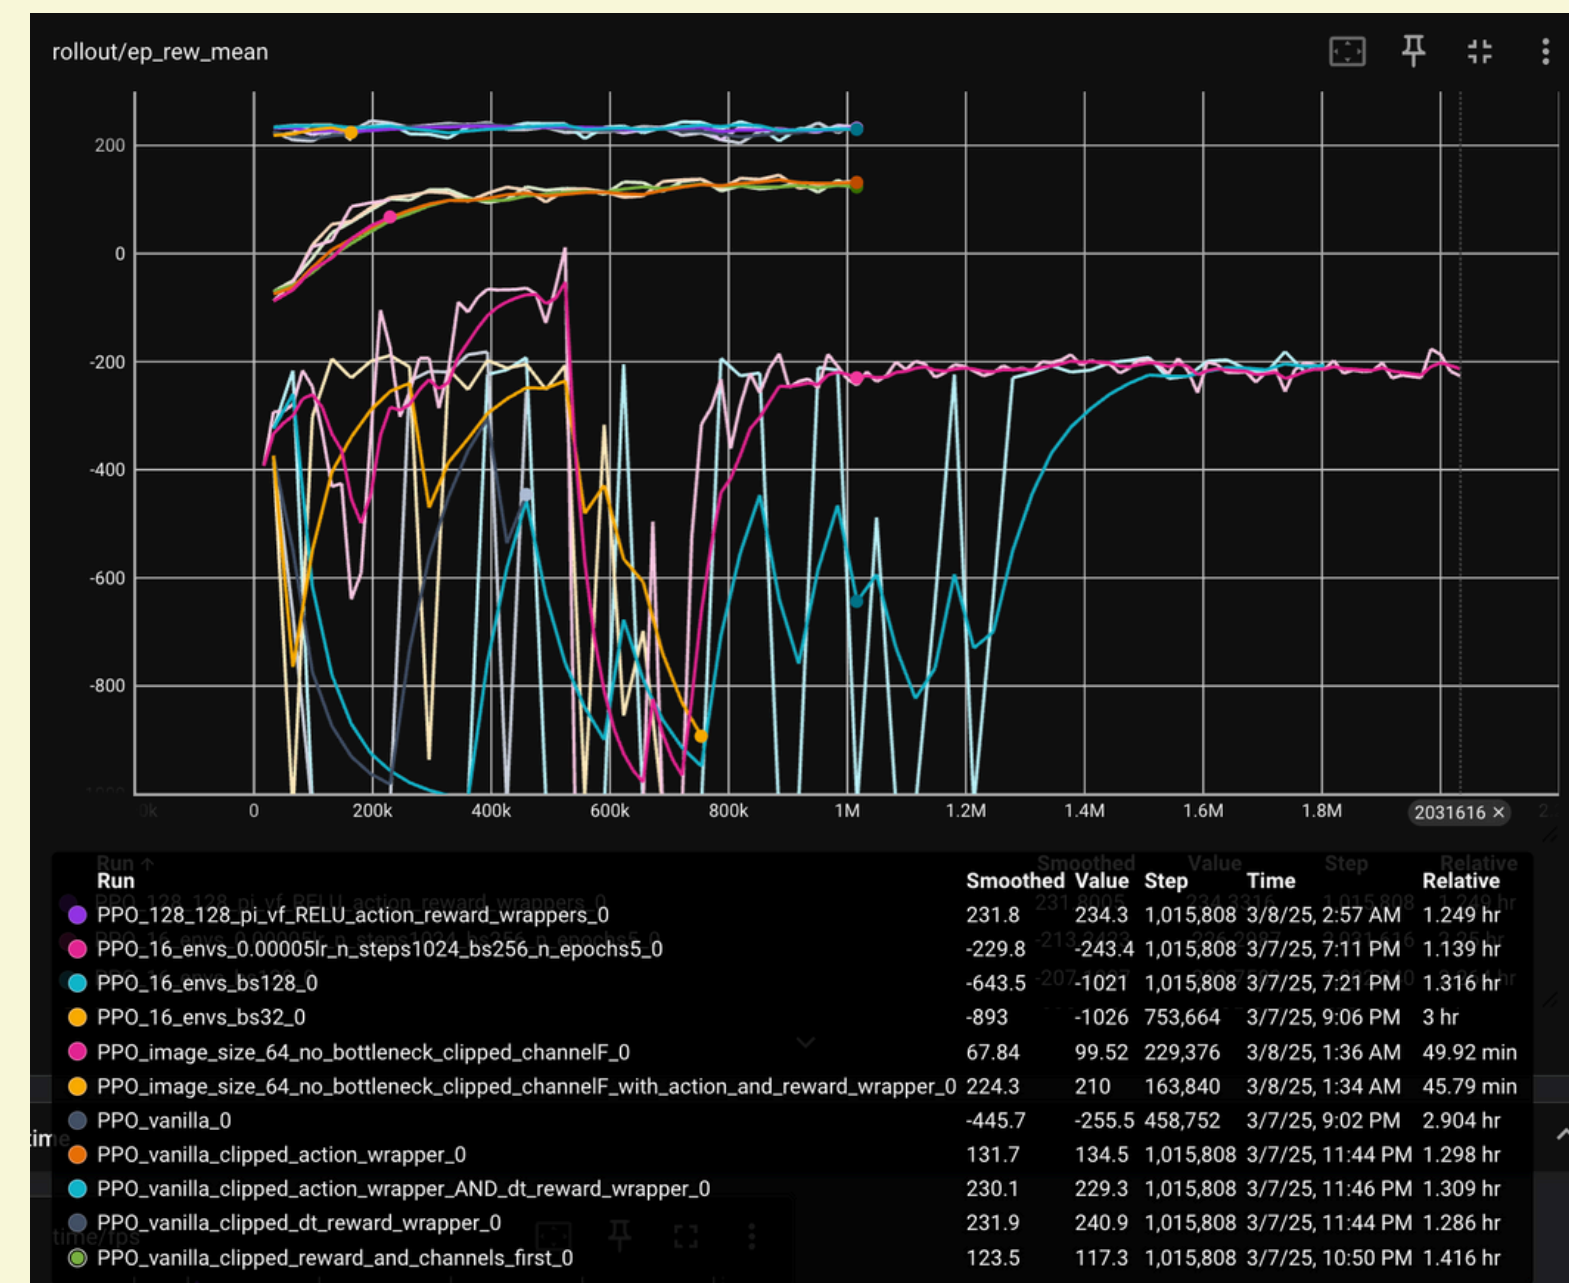
# PROXIMAL POLICY OPTIMIZATION (PPO)

- On-policy learning is preferred for its simplicity and stable updates
  - This is especially useful when using CNN inputs for navigation
  - CNN extracts lane and object features
  - Inputs predict actions like steering and velocity



# PPO TRAINING

- Built on Exercise 2 with additional **hyperparameter tuning** for optimization.
- Designed a **custom reward system** to improve learning.
- Trained for **5M timesteps** in both an **empty loop** and a **loop with obstacles**.
- Monitored progress via **TensorBoard** and recorded videos every **250K timesteps**.



# CHALLENGE #1: EXCESSIVE IMAGE SIZE AND SLOW TRAINING

## Issue:

- Native Duckietown camera resolution
  - **640x480**
- Large observation space resulted in extremely large models (~5GB).
- Training was **excessively slow**, with episodes lasting **too long** (high max\_steps).
- Debugging and experimentation became challenging.

## Solution:

- Fixed camera resolution
  - **3x84x84 (3 color channels)**
- **Significantly simplified** the observation space.
- Dramatically **decreased model size** and **memory usage**.
- **Training speed improved substantially**, enabling rapid experimentation.



# CHALLENGE #2: INEFFECTIVE DEFAULT REWARD FUNCTION

## Issue:

- Duckietown's default reward function provided **sparse** and **insufficient feedback**.
- **Poor guidance** for the agent in a continuous, dynamic environment.
- Agent **frequently failed** to stay within lanes or navigate smoothly.

## Solution:

- Designed a custom, dense reward function, providing immediate feedback.
- **Implemented multiple reward components:**
  - **Distance traveled along the track reward**
  - **Lane position, Proximity, and Collision penalties**

# HYPER PARAMETERS

- **Model Parameters**

- Learning Rate: 0.0001
- N Steps: 4096
- Batch Size: 256
- N Epochs: 10
- Gamma: 0.98
- GAE Lambda: 0.95
- Clip Range: 0.3
- Entropy Coefficient (ent\_coef): 0.05
- Value Function Coefficient (vf\_coef): 0.5
- Max Gradient Norm: 0.5

- **Simulator Parameters**

- Domain Randomization: False
- Max Steps per Episode: 250
- Draw Curve: False
- Draw Bounding Box: False
- Camera Width: 84
- Camera Height: 84

- **Training Configuration**

- Seed: 47
- Number of Environments: 8
- Reinforcement Learning Algorithm: PPO
- Total Timesteps: 5,000,000

# CUSTOM REWARD FUNCTION

## 1. Encouraging Forward Motion:

- The agent is rewarded for moving forward in a stable manner.
- Prevents the vehicle from spinning in place or getting stuck.

## 2. Penalizing Lane Deviations:

- Negative rewards are applied when the vehicle moves out of its lane.
- Encourages alignment with the center of the lane to improve navigation.

## 3. Minimizing Collisions:


- A heavy penalty is assigned when the vehicle collides with obstacles.
- Ensures the model learns to avoid objects and drive cautiously.

## 4. Discouraging Reverse Driving:

- A slight penalty is given when the agent moves in reverse unnecessarily.
- Reinforces forward driving as the optimal strategy.

## 5. Lane-Centric Corrections:

- The agent receives feedback based on lane-center distance and angle.
- Helps refine turning and smooth lane-following behavior.




# REWARD FUNCTION IMPROVEMENT: PREVENTING REVERSING BEHAVIOR

## Additional Observation:

- Early models often reversed to correct their mistakes.
- Frequent reversing resulted in erratic and inefficient driving patterns ("squiggly paths").
- This behavior hindered stable progression and reduced overall rewards.

## Solution:

- Introduced a severe penalty for reversing movements.
- Agent motivated explicitly to move forwards or sideways—never backward.
- Stabilizing driving results:
  - Smoother trajectories.
  - Consistent lane following.
  - Overall better agent performance and stability.




```
# Check if reversing (both left and right engine values are negative)
left_engine, right_engine = action
if left_engine < 0 and right_engine < 0:
    # Apply penalty for reversing
    return -25
```



# RESULTS

- Our agent can drive quickly, smoothly, and keep centering the lane in the small empty loop.
- Our agent can avoid obstacles in our custom map although inconsistent



The image features a light cream background with four yellow abstract shapes in the corners. The top-left shape consists of several thin, concentric, teardrop-like outlines. The top-right shape is a solid yellow, thick, wavy line. The bottom-left shape is a solid yellow, thick, wavy line. The bottom-right shape is a solid yellow, thick, wavy line.

**THANK  
YOU**